



(a) Method



$$L_{grad} = \| \nabla_u V(\hat{f}(x, u)) - \nabla_u V(f(x, u)) \|^2$$

(b) Logic



value_only fails everywhere → + L_{grad} rescues (not via contractivity); limitation: full sufficiency also needs prediction accuracy (axis a)